

# Randomized Experiment

```
library(conflicted)
```

```
library(kableExtra)
```

```
library(knitr)
```

```
library(broom.helpers)
```

```
library(broom)
```

```
library(dtplyr)
```

```
library(furrr)
```

```
## Loading required package: future
```

```
library(arrow)
```

```
library(glue)
```

```
library(fs)
```

```
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
```

```
## v dplyr      1.1.4      v readr      2.1.5
```

```
## v forcats    1.0.0      v stringr    1.5.1
```

```
## v ggplot2    3.5.1      v tibble     3.2.1
```

```
## v lubridate  1.9.3      v tidyr      1.3.1
```

```
## v purrr      1.0.2
```

```
conflict_prefer("filter", "dplyr")
```

```
## [conflicted] Will prefer dplyr::filter over any other package.
```

```
source(here("analysis/utils.R"), local = knitr_global())
```

```
set_theme()
```

```
write_bib(.packages(), here("analysis/packages.bib"))
```

```
sessionInfo()
```

```
## R version 4.4.0 (2024-04-24)
```

```
## Platform: aarch64-apple-darwin20
```

```
## Running under: macOS Sonoma 14.5
```

```
##
```

```
## Matrix products: default
```

```
## BLAS:   /Library/Frameworks/R.framework/Versions/4.4-arm64/Resources/lib/libRblas.0.dylib
```

```
## LAPACK: /Library/Frameworks/R.framework/Versions/4.4-arm64/Resources/lib/libRlapack.dylib; LAPACK v
```

```
##
```

```
## locale:
```

```
## [1] en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8
```

```
##
```

```
## time zone: Asia/Singapore
```

```
## tzcode source: internal
```

```
##
```

```
## attached base packages:
```

```
## [1] stats      graphics  grDevices  utils      datasets  methods   base
```

```
##
## other attached packages:
## [1] lubridate_1.9.3      forcats_1.0.0      stringr_1.5.1
## [4] dplyr_1.1.4          purrr_1.0.2        readr_2.1.5
## [7] tidyr_1.3.1          tibble_3.2.1       ggplot2_3.5.1
## [10] tidyverse_2.0.0      fs_1.6.4           glue_1.7.0
## [13] arrow_16.1.0         frr_0.3.1          future_1.33.2
## [16] dtplyr_1.3.1         broom_1.0.6        broom.helpers_1.15.0
## [19] knitr_1.47           kableExtra_1.4.0   conflicted_1.2.0
## [22] here_1.0.1
##
## loaded via a namespace (and not attached):
## [1] gtable_0.3.5         xfun_0.45          tzdb_0.4.0         vctrs_0.6.5
## [5] tools_4.4.0          generics_0.1.3     parallel_4.4.0     fansi_1.0.6
## [9] pkgconfig_2.0.3      data.table_1.15.4  assertthat_0.2.1   lifecycle_1.0.4
## [13] compiler_4.4.0       munsell_0.5.1      codetools_0.2-20   htmltools_0.5.8.1
## [17] yaml_2.3.8           pillar_1.9.0       cachem_1.1.0       parallelly_1.37.1
## [21] tidyselect_1.2.1     digest_0.6.35      stringi_1.8.4      listenv_0.9.1
## [25] rprojroot_2.0.4      fastmap_1.2.0      grid_4.4.0         colorspace_2.1-0
## [29] cli_3.6.2            magrittr_2.0.3     utf8_1.2.4         withr_3.0.0
## [33] scales_1.3.0         backports_1.5.0    bit64_4.0.5        timechange_0.3.0
## [37] rmarkdown_2.27       globals_0.16.3     bit_4.0.5          hms_1.1.3
## [41] memoise_2.0.1        evaluate_0.24.0    viridisLite_0.4.2  rlang_1.1.4
## [45] xml2_1.3.6           svglite_2.1.3      rstudioapi_0.16.0  R6_2.5.1
## [49] systemfonts_1.1.0
```

## Analyze attack trends

```
data_dir <- here(glue("{params$data}/{params$simulation}/results"))

success_fnames <-
  dir_ls(data_dir, glob = glue("*norm_{params$norm}*.csv"))

stopifnot(length(success_fnames) == 1200)

# every fname is a simulation
success_raw_data <- get_data(success_fnames, read_csv) |>
  glimpse()
```

```
## Rows: 1,200
## Columns: 16
## $ fname                <chr> "/Users/zbli/Documents/Documents - ZhaoBin's M-
## $ num_iteration         <dbl> 100, 100, 100, 100, 100, 100, 100, 100, 100, 1~
## $ max_norm              <dbl> 0.05, 0.05, 0.05, 0.05, 0.05, 0.05, 0.05, 0.05~
## $ model_name            <ord> Cascade R-CNN, Faster R-CNN, RetinaNet, SSD, Y~
## $ loss_target           <ord> Mislabeling, Mislabeling, Mislabeling, Mislabel~
## $ attack_bbox          <chr> "predictions", "predictions", "predictions", "~
## $ perturb_fun           <chr> "perturb_inside", "perturb_inside", "perturb_i~
## $ sample_count          <dbl> 247, 253, 258, 266, 261, 247, 253, 258, 266, 2~
## $ attack_count          <dbl> 200, 200, 200, 200, 200, 200, 200, 200, 200, 2~
## $ success_count         <dbl> 4, 6, 4, 31, 42, 7, 8, 9, 28, 13, 5, 4, 11, 40~
## $ vanish_count         <dbl> 2, 5, 0, 11, 14, 7, 8, 6, 22, 10, 5, 4, 10, 39~
## $ mislabel_count        <dbl> 2, 1, 4, 20, 28, 0, 0, 3, 6, 3, 0, 0, 1, 1, 0,~
```

```

## $ mislabel_intended_count <dbl> 2, 1, 4, 20, 27, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ~
## $ target_max_conf <lgl> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA~
## $ perturb_min_size <lgl> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA~
## $ bbox_max_dist <lgl> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA~

itr_lab <- "Attack Iterations"

cap <- glue("{emp_tex('Intent obfuscating attack is feasible for all models and attacks', params$norm)}")

cap

## Intent obfuscating attack is feasible for all models and attacks even with 0.05 max-norm: We conduct
success_intended_data <- success_raw_data |>
  mutate(success_intended_count = case_when(
    loss_target == "Mislabeling" ~ mislabel_intended_count,
    loss_target == "Vanishing" ~ vanish_count,
    loss_target == "Untargeted" ~ success_count
  ))

# expand intended success per simulation into 1 and 0s per row
success_expanded_data <- success_intended_data |>
  rowwise() |>
  mutate(success = list(rep(0:1, times = c(attack_count - success_intended_count, success_intended_count))))
  unnest_longer(success) |>
  glimpse()

## Rows: 240,000
## Columns: 18
## $ fname <chr> "/Users/zbli/Documents/Documents - ZhaoBin's M
## $ num_iteration <dbl> 100, 100, 100, 100, 100, 100, 100, 100, 100, 1~
## $ max_norm <dbl> 0.05, 0.05, 0.05, 0.05, 0.05, 0.05, 0.05, 0.05, 0.05~
## $ model_name <ord> Cascade R-CNN, Cascade R-CNN, Cascade R-CNN, C~
## $ loss_target <ord> Mislabeling, Mislabeling, Mislabeling, Mislabel~
## $ attack_bbox <chr> "predictions", "predictions", "predictions", "~
## $ perturb_fun <chr> "perturb_inside", "perturb_inside", "perturb_i~
## $ sample_count <dbl> 247, 247, 247, 247, 247, 247, 247, 247, 247, 2~
## $ attack_count <dbl> 200, 200, 200, 200, 200, 200, 200, 200, 200, 2~
## $ success_count <dbl> 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4~
## $ vanish_count <dbl> 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2~
## $ mislabel_count <dbl> 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2~
## $ mislabel_intended_count <dbl> 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2~
## $ target_max_conf <lgl> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA~
## $ perturb_min_size <lgl> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA~
## $ bbox_max_dist <lgl> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA~
## $ success_intended_count <dbl> 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2~
## $ success <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~

# use log(num_iteration)
g <- success_expanded_data |>
  ggplot(aes(num_iteration, success, color = loss_target, linetype = loss_target)) +
  # use stat_summary rather than stat_summary_bin
  # since num_iteration is set experimentally
  # mean_cl_boot gives 95% bootstrapped CI at 1000 samples
  # https://rdrr.io/cran/Hmisc/man/smean.sd.html
  stat_summary(fun.data = "mean_cl_boot") +

```




Figure 1: Intent obfuscating attack is feasible for all models and attacks even with 0.05 max-norm: We conduct a randomized experiment by resampling COCO images, and within those images randomly sampling correctly predicted target and perturb objects. Then we distort the perturb objects to disrupt the target objects varying the attack iterations. The binned summaries and regression trendlines graph success proportion against attack iterations in the randomized attack experiment. Errors are 95% confidence intervals and every point aggregates success over 4,000 images. Targeted vanishing and mislabeling attacks obtain significantly greater success on the 1-stage YOLOv3 and SSD than the 2-stage Faster R-CNN and Cascade R-CNN detectors. However, the 1-stage RetinaNet is as resilient as the 2-stage detectors. Moreover, success rates significantly increase with larger attack iterations. Significance is determined at  $\alpha < 0.05$  using a Wald z-test on the logistic estimates. Full details are given in Section ??.

```
binomial_smooth(formula = y ~ log(x)) +
facet_grid(cols = vars(model_name))

g +
  labs(x = itr_lab, y = glue("p(Success) {norm_axy(params$norm)}"), color = "Attack", linetype = "Attack"),
  scale_x_continuous(breaks = unique(success_raw_data$num_iteration))

# compare models against YOLO
# grouped by attack
data <- success_expanded_data |>
  # restrict to max iteration
  filter(num_iteration == max(num_iteration)) |>
  # avoid ordered regression
  mutate(
    model_name = factor(model_name, ordered = FALSE),
```

```

    loss_target = factor(loss_target, ordered = FALSE)
  ) |>
  glimpse()

## Rows: 60,000
## Columns: 18
## $ fname                <chr> "/Users/zbli/Documents/Documents - ZhaoBin's M-
## $ num_iteration        <dbl> 200, 200, 200, 200, 200, 200, 200, 200, 200, 2~
## $ max_norm             <dbl> 0.05, 0.05, 0.05, 0.05, 0.05, 0.05, 0.05, 0.05~
## $ model_name           <fct> Cascade R-CNN, Cascade R-CNN, Cascade R-CNN, C~
## $ loss_target          <fct> Mislabeling, Mislabeling, Mislabeling, Mislabel~
## $ attack_bbox         <chr> "predictions", "predictions", "predictions", "~
## $ perturb_fun          <chr> "perturb_inside", "perturb_inside", "perturb_i~
## $ sample_count         <dbl> 247, 247, 247, 247, 247, 247, 247, 247, 247, 2~
## $ attack_count        <dbl> 200, 200, 200, 200, 200, 200, 200, 200, 200, 2~
## $ success_count       <dbl> 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8~
## $ vanish_count       <dbl> 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2~
## $ mislabel_count      <dbl> 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6~
## $ mislabel_intended_count <dbl> 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6~
## $ target_max_conf     <lgl> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA~
## $ perturb_min_size    <lgl> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA~
## $ bbox_max_dist       <lgl> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA~
## $ success_intended_count <dbl> 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6~
## $ success             <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~

model <- partial(glm_model, predictor = "model_name")

reg_est <- get_tidied_reg(
  model, data, loss_target
)

## Warning: Returning more (or less) than 1 row per `summarise()` group was deprecated in
## dplyr 1.1.0.
## i Please use `reframe()` instead.
## i When switching from `summarise()` to `reframe()`, remember that `reframe()`
## always returns an ungrouped data frame and adjust accordingly.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.

## `summarise()` has grouped output by 'loss_target'. You can override using the
## `.groups` argument.

ext_sig(reg_est)

## Total 15 predictors:
## 10 (67%) significant;
## 10 (67%) both

## # A tibble: 10 x 8
## # Groups:   loss_target [3]
##   loss_target term      estimate std.error statistic p.value conf.low conf.high
##   <fct>      <chr>      <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
## 1 Vanishing  model_na~ -0.315    0.053    -5.96     0      -0.419   -0.211
## 2 Vanishing  model_na~ -1.72     0.075   -22.9     0      -1.88    -1.58
## 3 Vanishing  model_na~ -2.51     0.102   -24.7     0      -2.72    -2.32
## 4 Vanishing  model_na~ -1.95     0.082   -23.9     0      -2.12    -1.80

```

```
## 5 Mislabeling model_na~ -2.17      0.135    -16.1    0      -2.45    -1.92
## 6 Mislabeling model_na~ -2.94      0.189    -15.5    0      -3.33    -2.59
## 7 Mislabeling model_na~ -1.96      0.123    -15.9    0      -2.21    -1.72
## 8 Untargeted  model_na~  0.587     0.079     7.46    0       0.433    0.742
## 9 Untargeted  model_na~ -0.319    0.094    -3.39   0.001   -0.504   -0.135
## 10 Untargeted model_na~ -0.488    0.098    -4.95    0      -0.682   -0.296
```

```
cap <- table_caption("detection models, split by attack,", "Both vanishing and mislabeling attacks obtain higher success on 1-stage (YOLOv3, SSD) than 2-stage (Faster R-CNN, Cascade R-CNN) detectors. However, the 1-stage RetinaNet is as resilient as 2-stage detectors. Table headers are explained in Appendix ??.")
print_statistics(reg_est, cap)
```

Table 1: We run a logistic model regressing success against detection models, split by attack, in the randomized attack experiment. Both vanishing and mislabeling attacks obtain higher success on 1-stage (YOLOv3, SSD) than 2-stage (Faster R-CNN, Cascade R-CNN) detectors. However, the 1-stage RetinaNet is as resilient as 2-stage detectors. Table headers are explained in Appendix ??.

Group		Regression						
Attack	term	sig	estimate	std.error	statistic	p.value	conf.low	conf.high
Vanishing	YOLOv3		0.000					
	SSD	*	-0.315	0.053	-5.956	0.000	-0.419	-0.211
	RetinaNet	*	-1.725	0.075	-22.889	0.000	-1.875	-1.579
	Faster R-CNN	*	-2.511	0.102	-24.732	0.000	-2.715	-2.317
	Cascade R-CNN	*	-1.953	0.082	-23.914	0.000	-2.116	-1.796
Mislabeling	YOLOv3		0.000					
	SSD		-0.051	0.068	-0.751	0.453	-0.185	0.083
	RetinaNet	*	-2.173	0.135	-16.124	0.000	-2.446	-1.917
	Faster R-CNN	*	-2.939	0.189	-15.521	0.000	-3.332	-2.587
	Cascade R-CNN	*	-1.959	0.123	-15.888	0.000	-2.207	-1.723
Untargeted	YOLOv3		0.000					
	SSD	*	0.587	0.079	7.460	0.000	0.433	0.742
	RetinaNet		0.038	0.087	0.433	0.665	-0.132	0.208
	Faster R-CNN	*	-0.319	0.094	-3.389	0.001	-0.504	-0.135
	Cascade R-CNN	*	-0.488	0.098	-4.954	0.000	-0.682	-0.296

```
# compare attacks against vanishing
# grouped by models
model <- partial(glm_model, predictor = "loss_target")

reg_est <- get_tidied_reg(
  model, data, model_name
)
```

```
## Warning: Returning more (or less) than 1 row per `summarise()` group was deprecated in
## dplyr 1.1.0.
## i Please use `reframe()` instead.
## i When switching from `summarise()` to `reframe()`, remember that `reframe()`
## always returns an ungrouped data frame and adjust accordingly.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
```

```
## generated.

## `summarise()` has grouped output by 'model_name'. You can override using the
## `.groups` argument.

ext_sig(reg_est)

## Total 15 predictors:
## 9 (60%) significant;
## 9 (60%) both

## # A tibble: 9 x 8
## # Groups:   model_name [5]
##   model_name term      estimate std.error statistic p.value conf.low conf.high
##   <fct>      <chr>      <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
## 1 YOLOv3     loss_ta~ -0.928    0.06     -15.5     0        -1.05    -0.812
## 2 YOLOv3     loss_ta~ -1.56     0.071    -21.9     0        -1.70    -1.42
## 3 SSD        loss_ta~ -0.665    0.062    -10.7     0        -0.787   -0.543
## 4 SSD        loss_ta~ -0.66     0.062    -10.6     0        -0.783   -0.538
## 5 RetinaNet  loss_ta~ -1.38     0.142     -9.67     0        -1.66    -1.10
## 6 RetinaNet  loss_ta~  0.201    0.09      2.24     0.025     0.025    0.378
## 7 Faster R-CNN loss_ta~ -1.36     0.206     -6.57     0        -1.78    -0.966
## 8 Faster R-CNN loss_ta~  0.631    0.119      5.32     0         0.401    0.866
## 9 Cascade R-CNN loss_ta~ -0.934    0.135     -6.90     0        -1.20    -0.673

cap <- table_caption("attacks, split by detection models", "Targeted attacks obtain higher success than
print_statistics(reg_est, cap)
```

Table 2: We run a logistic model regressing success against attacks, split by detection models in the randomized attack experiment. Targeted attacks obtain higher success than untargeted attacks on YOLOv3 and SSD. Within targeted attacks, vanishing attacks obtain higher success than mislabeling attacks on all models. Table headers are explained in Appendix ??.

Group	Regression							
Model	term	sig	estimate	std.error	statistic	p.value	conf.low	conf.high
YOLOv3	Vanishing		0.000					
	Mislabeling	*	-0.928	0.060	-15.542	0.000	-1.046	-0.812
	Untargeted	*	-1.561	0.071	-21.871	0.000	-1.703	-1.423
SSD	Vanishing		0.000					
	Mislabeling	*	-0.665	0.062	-10.658	0.000	-0.787	-0.543
	Untargeted	*	-0.660	0.062	-10.594	0.000	-0.783	-0.538
RetinaNet	Vanishing		0.000					
	Mislabeling	*	-1.376	0.142	-9.667	0.000	-1.663	-1.104
	Untargeted	*	0.201	0.090	2.237	0.025	0.025	0.378
Faster R-CNN	Vanishing		0.000					
	Mislabeling	*	-1.356	0.206	-6.571	0.000	-1.778	-0.966
	Untargeted	*	0.631	0.119	5.317	0.000	0.401	0.866
	Vanishing		0.000					
	Mislabeling	*	-0.934	0.135	-6.901	0.000	-1.204	-0.673

Cascade R-CNN							
	Untargeted	-0.096	0.106	-0.901	0.367	-0.304	0.112

```
# num_iteration
reg_est <- get_tidied_reg(
  partial(glm_model, predictor = "log(num_iteration)"),
  success_expanded_data,
)

## `summarise()` has grouped output by 'model_name', 'loss_target'. You can
## override using the `.groups` argument.
ext_sig(reg_est, "pos")

## Total 15 predictors:
## 15 (100%) significant;
## 15 (100%) pos

## # A tibble: 15 x 9
## # Groups:   model_name, loss_target [15]
##   model_name loss_target term estimate std.error statistic p.value conf.low
##   <ord>      <ord>      <chr>   <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
## 1 YOLOv3     Vanishing log(~)  0.797    0.027    29.7      0    0.745
## 2 YOLOv3     Mislabeling log(~)  1.10     0.051    21.6      0     1
## 3 YOLOv3     Untargeted log(~)  0.347    0.036     9.62     0    0.277
## 4 SSD        Vanishing log(~)  0.852    0.032    26.6      0    0.79
## 5 SSD        Mislabeling log(~)  0.922    0.044    20.9      0    0.837
## 6 SSD        Untargeted log(~)  0.483    0.031    15.7      0    0.423
## 7 RetinaNet  Vanishing log(~)  0.88     0.062    14.2      0    0.762
## 8 RetinaNet  Mislabeling log(~)  0.903    0.115     7.86     0    0.688
## 9 RetinaNet  Untargeted log(~)  0.627    0.046    13.6      0    0.538
## 10 Faster R-CNN Vanishing log(~)  0.707    0.082     8.66     0    0.552
## 11 Faster R-CNN Mislabeling log(~)  0.975    0.191     5.11     0    0.627
## 12 Faster R-CNN Untargeted log(~)  0.483    0.049     9.94     0    0.389
## 13 Cascade R-CNN Vanishing log(~)  0.738    0.062    11.8      0    0.619
## 14 Cascade R-CNN Mislabeling log(~)  1.25     0.149     8.40     0    0.972
## 15 Cascade R-CNN Untargeted log(~)  0.45     0.05     9.04     0    0.354
## # i 1 more variable: conf.high <dbl>

cap <- table_caption(glue("log({itr_lab})"), "Success rates increase with attack iterations for all models")

print_statistics(reg_est, cap)
```

Table 3: We run a logistic model regressing success against log(attack iterations) in the randomized attack experiment. Success rates increase with attack iterations for all models and attacks. Table headers are explained in Appendix ??.

Group		Regression						
Attack	term	sig	estimate	std.error	statistic	p.value	conf.low	conf.high
<b>YOLOv3</b>								
Vanishing	log(iterations)	*	0.797	0.027	29.736	0	0.745	0.850
Mislabeling	log(iterations)	*	1.097	0.051	21.572	0	1.000	1.199
Untargeted	log(iterations)	*	0.347	0.036	9.615	0	0.277	0.419



<b>SSD</b>									
Vanishing	log(iterations)	*	0.852	0.032	26.573	0	0.790	0.915	
Mislabeling	log(iterations)	*	0.922	0.044	20.885	0	0.837	1.010	
Untargeted	log(iterations)	*	0.483	0.031	15.652	0	0.423	0.544	
<b>RetinaNet</b>									
Vanishing	log(iterations)	*	0.880	0.062	14.229	0	0.762	1.005	
Mislabeling	log(iterations)	*	0.903	0.115	7.855	0	0.688	1.139	
Untargeted	log(iterations)	*	0.627	0.046	13.591	0	0.538	0.719	
<b>Faster R-CNN</b>									
Vanishing	log(iterations)	*	0.707	0.082	8.664	0	0.552	0.872	
Mislabeling	log(iterations)	*	0.975	0.191	5.111	0	0.627	1.378	
Untargeted	log(iterations)	*	0.483	0.049	9.938	0	0.389	0.580	
<b>Cascade R-CNN</b>									
Vanishing	log(iterations)	*	0.738	0.062	11.832	0	0.619	0.863	
Mislabeling	log(iterations)	*	1.248	0.149	8.395	0	0.972	1.556	
Untargeted	log(iterations)	*	0.450	0.050	9.040	0	0.354	0.549	

## Analyze individual cases

```
# cache.lazy = FALSE needed to avoid errors with large bbox .parquets
attack_bbox <- "predictions"
```

```
bbox_fnames <-
  dir_ls(data_dir, glob = glue("*{params$norm}*.parquet"))
```

```
# Every bbox whether ground-truth, predicted or attacked is a row and the columns are the sample and bb
bbox_raw_data <- get_data(bbox_fnames, combine_trend_case) |>
  glimpse() |>
  lazy_dt()
```

```
## Rows: 8,712,402
## Columns: 41
## $ fname                <chr> "/Users/zbli/Documents/Documents - ZhaoBin"~
## $ sample_id            <chr> "65ed3a88141a475067f32706", "65ed3a88141a47~
## $ sample_path          <chr> "/projects/f_ps848_1/zhaobin/adversarial/co~
## $ sample_width         <int> 640, 640, 640, 640, 640, 640, 640, 640, 640~
## $ sample_height        <int> 480, 480, 480, 480, 480, 480, 480, 480, 480~
## $ sample_mislabel_class <chr> "horse", "horse", "horse", "horse", "horse"~
## $ sample_mislabel_proba <dbl> 6.615031e-05, 6.615031e-05, 6.615031e-05, 6~
## $ sample_attack        <lgl> TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, T~
## $ sample_vanish        <lgl> FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, F~
## $ sample_mislabel_intended <lgl> FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, F~
## $ sample_success        <lgl> FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, F~
## $ sample_mislabel       <lgl> FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, F~
## $ bbox_id              <chr> "65ed3a88141a475067f32700", "65ed3a88141a47~
## $ bbox_class            <chr> "clock", "person", "person", "person", "per~
## $ bbox_xywhn            <list<double>> <0.32484375, 0.26458333, 0.0474218~
## $ bbox_conf             <dbl> NA, NA, NA, NA, NA, NA, 0.9890913, 0.986363~
## $ bbox_res_eval         <chr> "tp", "tp", "tp", "tp", "tp", "fn", "tp", "~
```

```
## $ bbox_iou_eval <dbl> 0.8860679, 0.8505562, 0.8757091, 0.8901640,~
## $ bbox_res_pgd_eval <chr> NA, NA, NA, NA, NA, NA, "tp", "tp", "tp", "~
## $ bbox_iou_pgd_eval <dbl> NA, NA, NA, NA, NA, NA, 0.9999464, 0.999894~
## $ bbox_res_pgd_mislabel_eval <chr> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ bbox_iou_pgd_mislabel_eval <dbl> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ bbox_target <lg1> TRUE, FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, FA~
## $ bbox_perturb <lg1> FALSE, FALSE, FALSE, FALSE, TRUE, FALSE, FA~
## $ bbox_type <chr> "ground_truth", "ground_truth", "ground_tru~
## $ bbox_mislabel <lg1> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ num_iteration <dbl> 100, 100, 100, 100, 100, 100, 100, 100, 100, 100~
## $ max_norm <dbl> 0.05, 0.05, 0.05, 0.05, 0.05, 0.05, 0.05, 0~
## $ model_name <ord> Cascade R-CNN, Cascade R-CNN, Cascade R-CNN~
## $ loss_target <ord> Mislabeling, Mislabeling, Mislabeling, Misl~
## $ attack_bbox <chr> "predictions", "predictions", "predictions"~
## $ perturb_fun <chr> "perturb_inside", "perturb_inside", "pertur~
## $ sample_count <dbl> 247, 247, 247, 247, 247, 247, 247, 247, 247, 247~
## $ attack_count <dbl> 200, 200, 200, 200, 200, 200, 200, 200, 200, 200~
## $ success_count <dbl> 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4~
## $ vanish_count <dbl> 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2~
## $ mislabel_count <dbl> 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2~
## $ mislabel_intended_count <dbl> 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2~
## $ target_max_conf <lg1> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ perturb_min_size <lg1> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ bbox_max_dist <lg1> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,~
```

```
# check whether target and perturb bboxes and
# mislabel classes are seeded across iterations
cols_start_equal(bbox_raw_data, c(
  "bbox_target", "bbox_perturb",
  "sample_mislabel_class", "sample_mislabel_proba"
))
```

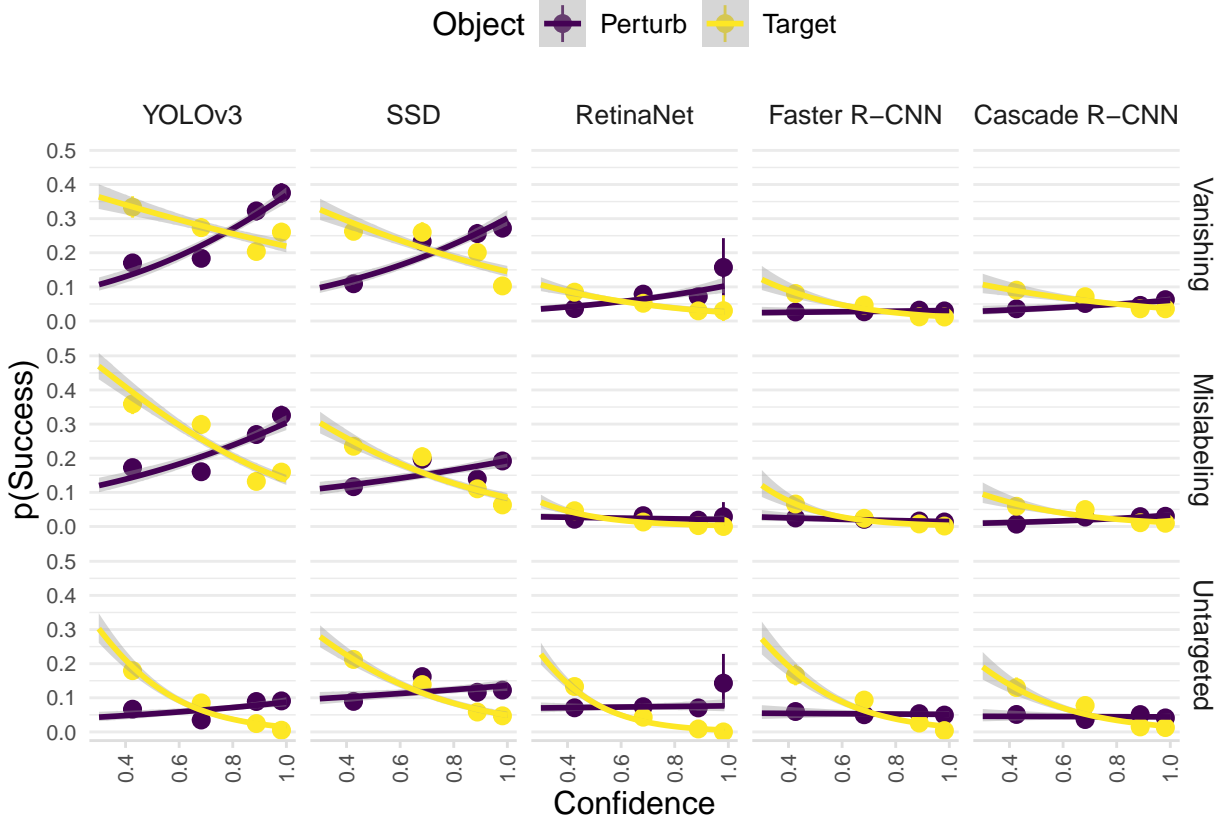
```
## Columns starting with `bbox_target` are equal: TRUE
## Columns starting with `bbox_perturb` are equal: TRUE
## Columns starting with `sample_mislabel_class` are equal: TRUE
## Columns starting with `sample_mislabel_proba` are equal: TRUE
```

```
# bbox confidence always based on predicted bbox
bbox_conf_data <- bbox_raw_data |>
  filter(bbox_type == "predictions") |>
  wrangle_success() |>
  glimpse()
```

```
## Rows: 120,000
## Columns: 42
## $ fname <chr> "/Users/zbli/Documents/Documents - ZhaoBin~
## $ sample_id <chr> "65ed3a88141a475067f32706", "65ed3a88141a47~
## $ sample_path <chr> "/projects/f_ps848_1/zhaobin/adversarial/co~
## $ sample_width <int> 640, 640, 500, 640, 480, 640, 640, 640, 640~
## $ sample_height <int> 480, 427, 332, 425, 640, 480, 480, 480, 640~
## $ sample_mislabel_class <chr> "horse", "truck", "surfboard", "horse", "ca~
## $ sample_mislabel_proba <dbl> 6.615031e-05, 4.219168e-02, 4.392489e-05, 1~
## $ sample_attack <lg1> TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, T~
## $ sample_vanish <lg1> FALSE, FALSE, TRUE, FALSE, FALSE, FALSE, FA~
## $ sample_mislabel_intended <lg1> FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, F~
## $ sample_success <lg1> FALSE, FALSE, TRUE, FALSE, FALSE, FALSE, FA~
```

```
## $ sample_mislabel      <lg1> FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, F~
## $ bbox_id              <chr> "65ed3aa3141a475067f3ca3e", "65ed3aa3141a47~
## $ bbox_class           <chr> "clock", "car", "person", "person", "donut"~
## $ bbox_xywhn           <list<double>> <0.32723613, 0.26601949, 0.0435188~
## $ bbox_conf            <dbl> 0.9305881, 0.3433506, 0.9882318, 0.9988949,~
## $ bbox_res_eval        <chr> "tp", "tp", "tp", "tp", "tp", "tp", "tp", "~
## $ bbox_iou_eval        <dbl> 0.8860679, 0.7609860, 0.9454082, 0.9299325,~
## $ bbox_res_pgd_eval    <chr> "tp", "tp", "fn", "tp", "tp", "tp", "tp", "~
## $ bbox_iou_pgd_eval    <dbl> 1.0000000, 1.0000000, NA, 0.9999969, 1.0000~
## $ bbox_res_pgd_mislabel_eval <chr> NA, NA, "fn", NA, NA, NA, NA, NA, NA, NA, N~
## $ bbox_iou_pgd_mislabel_eval <dbl> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ bbox_type            <chr> "predictions", "predictions", "predictions"~
## $ bbox_mislabel        <lg1> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ num_iteration        <dbl> 200, 200, 200, 200, 200, 200, 200, 200, 200, 200~
## $ max_norm             <dbl> 0.05, 0.05, 0.05, 0.05, 0.05, 0.05, 0.05, 0~
## $ model_name           <ord> Cascade R-CNN, Cascade R-CNN, Cascade R-CNN~
## $ loss_target          <ord> Mislabeling, Mislabeling, Mislabeling, Misl~
## $ attack_bbox          <chr> "predictions", "predictions", "predictions"~
## $ perturb_fun          <chr> "perturb_inside", "perturb_inside", "pertur~
## $ sample_count         <dbl> 247, 247, 247, 247, 247, 247, 247, 247, 247, 247~
## $ attack_count         <dbl> 200, 200, 200, 200, 200, 200, 200, 200, 200, 200~
## $ success_count        <dbl> 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8~
## $ vanish_count        <dbl> 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2~
## $ mislabel_count       <dbl> 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6~
## $ mislabel_intended_count <dbl> 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6~
## $ target_max_conf      <lg1> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ perturb_min_size     <lg1> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ bbox_max_dist        <lg1> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ target_or_perturb    <ord> Target, Target, Target, Target, Target, Target, Tar~
## $ target_or_perturb_boolean <lg1> TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, T~
## $ success              <dbl> 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
```

```
bbox_conf_data |>
  graph_attr(bbox_conf, "Confidence")
```



```
# restrict to target
pred_name <- "target confidence"
main_pt <- glue("Lower {pred_name} significantly increases success rates for all models and attacks")

bbox_conf_graph <- bbox_conf_data |> filter(target_or_perturb == "Target")
bbox_conf_graph |>
  graph_attr(bbox_conf, pred_name)

model <- partial(glm_model, predictor = "bbox_conf")
data <- bbox_conf_graph

reg_est <- get_tidied_reg(model, data)

## `summarise()` has grouped output by 'model_name', 'loss_target'. You can
## override using the `.groups` argument.

ext_sig(reg_est, "neg")
```

```
## Total 15 predictors:
## 15 (100%) significant;
## 15 (100%) neg

## # A tibble: 15 x 9
## # Groups:   model_name, loss_target [15]
##   model_name loss_target term estimate std.error statistic p.value conf.low
##   <ord>      <ord>      <chr>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
## 1 YOLOv3     Vanishing  bbox~   -1.02    0.162    -6.29     0    -1.33
## 2 YOLOv3     Mislabeling bbox~   -2.47    0.171   -14.4     0    -2.81
## 3 YOLOv3     Untargeted  bbox~   -4.84    0.313   -15.5     0    -5.47
```




Figure 2: Lower target confidence significantly increases success rates for all models and attacks even with 0.05 max-norm: The binned summaries and regression trendlines graph success proportion against target confidence in the randomized attack experiment. Bins are split into quantiles. Errors are 95% confidence intervals

```
## 4 SSD Vanishing bbox~ -1.50 0.163 -9.25 0 -1.82
## 5 SSD Mislabeling bbox~ -2.21 0.185 -12.0 0 -2.58
## 6 SSD Untargeted bbox~ -2.89 0.215 -13.5 0 -3.31
## 7 RetinaNet Vanishing bbox~ -2.20 0.36 -6.12 0 -2.92
## 8 RetinaNet Mislabeling bbox~ -4.78 0.682 -7.00 0 -6.17
## 9 RetinaNet Untargeted bbox~ -5.82 0.439 -13.2 0 -6.70
## 10 Faster R-CNN Vanishing bbox~ -3.44 0.39 -8.81 0 -4.21
## 11 Faster R-CNN Mislabeling bbox~ -5.24 0.56 -9.36 0 -6.38
## 12 Faster R-CNN Untargeted bbox~ -4.52 0.313 -14.4 0 -5.14
## 13 Cascade R-CNN Vanishing bbox~ -1.65 0.303 -5.43 0 -2.24
## 14 Cascade R-CNN Mislabeling bbox~ -3.15 0.412 -7.64 0 -3.96
## 15 Cascade R-CNN Untargeted bbox~ -3.81 0.326 -11.7 0 -4.46
## # i 1 more variable: conf.high <dbl>
```

```
print_statistics(reg_est, table_caption(pred_name, main_pt))
```

Table 4: We run a logistic model regressing success against target confidence in the randomized attack experiment. Lower target confidence significantly increases success rates for all models and attacks. Table headers are explained in Appendix ??.

Group	Regression
-------	------------

Attack	term	sig	estimate	std.error	statistic	p.value	conf.low	conf.high
<b>YOLOv3</b>								
Vanishing	confidence	*	-1.017	0.162	-6.286	0	-1.334	-0.700
Mislabeling	confidence	*	-2.470	0.171	-14.445	0	-2.806	-2.136
Untargeted	confidence	*	-4.845	0.313	-15.476	0	-5.470	-4.241
<b>SSD</b>								
Vanishing	confidence	*	-1.505	0.163	-9.251	0	-1.825	-1.187
Mislabeling	confidence	*	-2.212	0.185	-11.970	0	-2.576	-1.852
Untargeted	confidence	*	-2.889	0.215	-13.462	0	-3.313	-2.471
<b>RetinaNet</b>								
Vanishing	confidence	*	-2.203	0.360	-6.124	0	-2.918	-1.507
Mislabeling	confidence	*	-4.778	0.682	-7.002	0	-6.173	-3.491
Untargeted	confidence	*	-5.816	0.439	-13.241	0	-6.701	-4.977
<b>Faster R-CNN</b>								
Vanishing	confidence	*	-3.442	0.390	-8.814	0	-4.213	-2.680
Mislabeling	confidence	*	-5.244	0.560	-9.361	0	-6.383	-4.178
Untargeted	confidence	*	-4.522	0.313	-14.433	0	-5.144	-3.915
<b>Cascade R-CNN</b>								
Vanishing	confidence	*	-1.647	0.303	-5.433	0	-2.237	-1.047
Mislabeling	confidence	*	-3.146	0.412	-7.635	0	-3.960	-2.341
Untargeted	confidence	*	-3.811	0.326	-11.692	0	-4.456	-3.177

```
perturb_error_data <- bbox_conf_data |>
  filter(target_or_perturb == "Perturb") |>
  group_by(model_name, loss_target) |>
  summarise(perturb_error = 1 - mean(success)) |>
  glimpse()
```

## `summarise()` has grouped output by 'model\_name'. You can override using the  
## `.groups` argument.

```
## Rows: 15
## Columns: 3
## Groups: model_name [5]
## $ model_name      <ord> YOLOv3, YOLOv3, YOLOv3, SSD, SSD, SSD, RetinaNet, Retina~
## $ loss_target     <ord> Vanishing, Mislabeling, Untargeted, Vanishing, Mislabeli~
## $ perturb_error   <dbl> 0.73450, 0.76650, 0.92950, 0.79000, 0.84200, 0.88000, 0.~
```

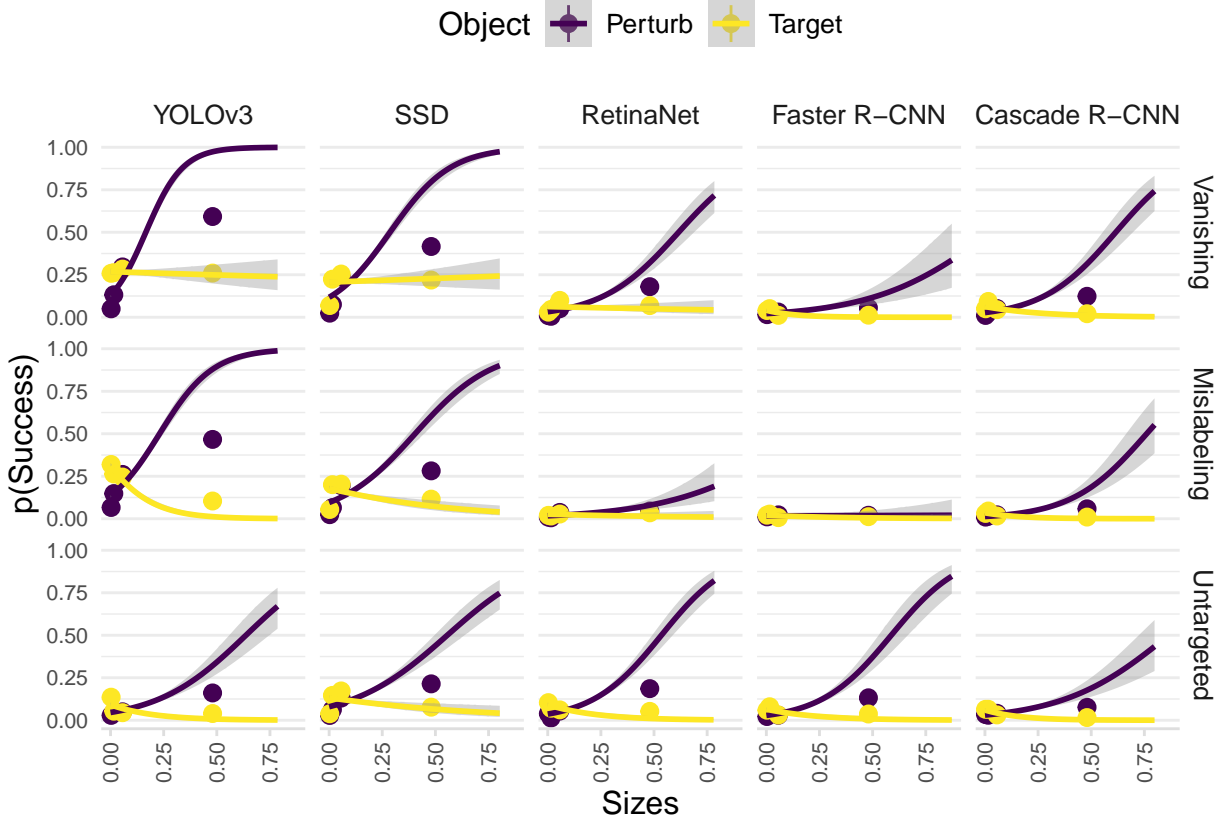
*# bbox sizes typically based on ground-truth attacked bbox*

```
bbox_size_data <- bbox_raw_data |>
  filter(bbox_type == attack_bbox) |>
  wrangle_success() |>
  # hoist not implemented in dplyr
  as_tibble() |>
  # bbox_xywhn == normalized x1, y1, w, h
  hoist(bbox_xywhn, bbox_xn = 1, bbox_yn = 2, bbox_wn = 3, bbox_hn = 4) |>
  mutate(
    bbox_size = bbox_wn * bbox_hn,
  ) |>
```

```
glimpse()
```

```
## Rows: 120,000
## Columns: 46
## $ fname <chr> "/Users/zbli/Documents/Documents - ZhaoBin'~
## $ sample_id <chr> "65ed3a88141a475067f32706", "65ed3a88141a47~
## $ sample_path <chr> "/projects/f_ps848_1/zhaobin/adversarial/co~
## $ sample_width <int> 640, 640, 500, 640, 480, 640, 640, 640, 640~
## $ sample_height <int> 480, 427, 332, 425, 640, 480, 480, 480, 640~
## $ sample_mislabel_class <chr> "horse", "truck", "surfboard", "horse", "ca~
## $ sample_mislabel_proba <dbl> 6.615031e-05, 4.219168e-02, 4.392489e-05, 1~
## $ sample_attack <lg1> TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, T~
## $ sample_vanish <lg1> FALSE, FALSE, TRUE, FALSE, FALSE, FALSE, FA~
## $ sample_mislabel_intended <lg1> FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, F~
## $ sample_success <lg1> FALSE, FALSE, TRUE, FALSE, FALSE, FALSE, FA~
## $ sample_mislabel <lg1> FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, F~
## $ bbox_id <chr> "65ed3aa3141a475067f3ca3e", "65ed3aa3141a47~
## $ bbox_class <chr> "clock", "car", "person", "person", "donut"~
## $ bbox_xn <dbl> 0.32723613, 0.81016169, 0.37364487, 0.58023~
## $ bbox_yn <dbl> 0.26601949, 0.50290289, 0.31231453, 0.46766~
## $ bbox_wn <dbl> 0.04351888, 0.03631706, 0.35480569, 0.08531~
## $ bbox_hn <dbl> 0.10756386, 0.02172394, 0.67813552, 0.40265~
## $ bbox_conf <dbl> 0.9305881, 0.3433506, 0.9882318, 0.9988949,~
## $ bbox_res_eval <chr> "tp", "tp", "tp", "tp", "tp", "tp", "tp", "~
## $ bbox_iou_eval <dbl> 0.8860679, 0.7609860, 0.9454082, 0.9299325,~
## $ bbox_res_pgd_eval <chr> "tp", "tp", "fn", "tp", "tp", "tp", "tp", "~
## $ bbox_iou_pgd_eval <dbl> 1.0000000, 1.0000000, NA, 0.9999969, 1.0000~
## $ bbox_res_pgd_mislabel_eval <chr> NA, NA, "fn", NA, NA, NA, NA, NA, NA, N~
## $ bbox_iou_pgd_mislabel_eval <dbl> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ bbox_type <chr> "predictions", "predictions", "predictions"~
## $ bbox_mislabel <lg1> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ num_iteration <dbl> 200, 200, 200, 200, 200, 200, 200, 200, 200~
## $ max_norm <dbl> 0.05, 0.05, 0.05, 0.05, 0.05, 0.05, 0.05, 0~
## $ model_name <ord> Cascade R-CNN, Cascade R-CNN, Cascade R-CNN~
## $ loss_target <ord> Mislabeling, Mislabeling, Mislabeling, Misl~
## $ attack_bbox <chr> "predictions", "predictions", "predictions"~
## $ perturb_fun <chr> "perturb_inside", "perturb_inside", "pertur~
## $ sample_count <dbl> 247, 247, 247, 247, 247, 247, 247, 247, 247~
## $ attack_count <dbl> 200, 200, 200, 200, 200, 200, 200, 200, 200~
## $ success_count <dbl> 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8~
## $ vanish_count <dbl> 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2~
## $ mislabel_count <dbl> 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6~
## $ mislabel_intended_count <dbl> 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6~
## $ target_max_conf <lg1> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ perturb_min_size <lg1> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ bbox_max_dist <lg1> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ target_or_perturb <ord> Target, Target, Target, Target, Target, Tar~
## $ target_or_perturb_boolean <lg1> TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, T~
## $ success <dbl> 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
## $ bbox_size <dbl> 0.0046810584, 0.0007889497, 0.2406063427, 0~
```

```
bbox_size_data |>
  graph_attr(bbox_size, "Sizes")
```



```
# bbox distances typically based on ground-truth attacked bbox as in sizes
bbox_dist_data <- bbox_size_data |>
mutate(
  target_or_perturb_lower = str_to_lower(target_or_perturb)
) |>
# mainly "group" by sample_id and attack iteration
# with target bbox on one row and perturb on another
# success, model_name, loss_target are sample attributes
# duplicated across bboxes
pivot_wider(
  id_cols = c(fname, sample_id, num_iteration, success, model_name, loss_target), names_from = target,
  values_from = c(bbox_xn, bbox_yn, bbox_wn, bbox_hn, bbox_size)
) |>
rowwise() |>
mutate(bbox_dist = get_min_distance(
  bbox_xn_perturb, bbox_yn_perturb, bbox_xn_perturb + bbox_wn_perturb, bbox_yn_perturb + bbox_hn_perturb,
  bbox_xn_target, bbox_yn_target, bbox_xn_target + bbox_wn_target, bbox_yn_target + bbox_hn_target
)) |>
ungroup() |>
glimpse()
```

```
## Rows: 60,000
```

```
## Columns: 17
```

```
## $ fname
```

```
## $ sample_id
```

```
## $ num_iteration
```

```
## $ success
```

```
## $ model_name
```

```
<chr> "/Users/zbli/Documents/Documents - ZhaoBin's MacBook~
```

```
<chr> "65ed3a88141a475067f32706", "65ed3a88141a475067f3272~
```

```
<dbl> 200, 200, 200, 200, 200, 200, 200, 200, 200, 200, 20~
```

```
<dbl> 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
```


```
<ord> Cascade R-CNN, Cascade R-CNN, Cascade R-CNN, Cascade~
```



```
## $ loss_target      <ord> Mislabeling, Mislabeling, Mislabeling, Mislabeling, ~
## $ bbox_xn_target   <dbl> 0.32723613, 0.81016169, 0.37364487, 0.58023462, 0.82~
## $ bbox_xn_perturb  <dbl> 4.478896e-01, 1.517359e-01, 3.132355e-02, 2.802266e-~
## $ bbox_yn_target   <dbl> 0.26601949, 0.50290289, 0.31231453, 0.46766415, 0.18~
## $ bbox_yn_perturb  <dbl> 0.8013828, 0.5229044, 0.7769909, 0.4782841, 0.469777~
## $ bbox_wn_target   <dbl> 0.04351888, 0.03631706, 0.35480569, 0.08531094, 0.07~
## $ bbox_wn_perturb  <dbl> 0.02720404, 0.07043431, 0.18098172, 0.13681064, 0.12~
## $ bbox_hn_target   <dbl> 0.10756386, 0.02172394, 0.67813552, 0.40265309, 0.04~
## $ bbox_hn_perturb  <dbl> 0.07742354, 0.04831026, 0.21702971, 0.48990981, 0.04~
## $ bbox_size_target <dbl> 0.0046810584, 0.0007889497, 0.2406063427, 0.03435071~
## $ bbox_size_perturb <dbl> 0.0021062328, 0.0034026994, 0.0392784101, 0.06702487~
## $ bbox_dist        <dbl> 0.43469769, 0.58799145, 0.16133960, 0.16319737, 0.28~
```

```
bbox_dist_data |>
  graph_attr(bbox_dist, "Distances")
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```



```
saveRDS(bbox_dist_data, here(glue("analysis/rand_dist_size_norm_{params$norm}.RDS")))
```

```
check_graph_data(bbox_dist_data, c(bbox_dist, bbox_size_perturb))
```

```
dist_lab <- "Perturb-Target Distance (relative to image width/height)"
```

```
size_lab <- "Perturb Box Size (relative to image width/height)"
```

```
pred_name <- glue("{dist_lab} and {size_lab}")
```

```
main_pt <- "Larger perturb objects significantly increase success rates for all models and attacks, exc"
```

```

cap <- glue(
  "{emp_tex(main_pt, params$norm)} The binned summaries",
  " graph success proportion against {str_to_lower(pred_name)} in the",
  " randomized attack experiment."
)

bbox_dist_data <- bbox_dist_data |> mutate(
  bbox_size_perturb = bbox_size_perturb,
  bbox_dist = bbox_dist
)

graph_dist_size <- function(g) {
  g + facet_grid(rows = vars(loss_target), cols = vars(model_name)) +
    labs(x = dist_lab, y = size_lab) +
    scale_fill_viridis_c(name = "p(Success)", breaks = c(0, .5, 1), limits = c(0, 1))
}

g <- bbox_dist_data |> ggplot(aes(bbox_dist, bbox_size_perturb, z = success)) +
  stat_summary_2d(fun = "mean", bins = 5)

graph_dist_size(g)

```

```

# control both
model <- partial(glm_model, predictor = "bbox_dist * bbox_size_perturb")
data <- bbox_dist_data

reg_res <- get_tidied_reg(model, data, return_mod = TRUE) |> glimpse()

```

```

## Warning: There were 4 warnings in `mutate()`.
## The first warning was:
## i In argument: `mod = list(model(data))`.
## i In row 7.
## Caused by warning:
## ! glm.fit: fitted probabilities numerically 0 or 1 occurred
## i Run `dplyr::last_dplyr_warnings()` to see the 3 remaining warnings.

## Warning: There were 168 warnings in `summarize()`.
## The first warning was:
## i In argument: `tidy_plus_plus(mod, conf.int = TRUE)`.
## i In row 7.
## Caused by warning:
## ! glm.fit: fitted probabilities numerically 0 or 1 occurred
## i Run `dplyr::last_dplyr_warnings()` to see the 167 remaining warnings.

## Warning: Returning more (or less) than 1 row per `summarise()` group was deprecated in
## dplyr 1.1.0.
## i Please use `reframe()` instead.
## i When switching from `summarise()` to `reframe()`, remember that `reframe()`
## always returns an ungrouped data frame and adjust accordingly.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.

## `summarise()` has grouped output by 'model_name', 'loss_target'. You can
## override using the `.groups` argument.

```




Figure 3: Larger perturb objects significantly increase success rates for all models and attacks, except for mislabeling attack on Faster R-CNN, after controlling for perturb-target distances. Shorter perturb-target distances significantly increase success rates for all models and attacks, after controlling for perturb object sizes even with 0.05 max-norm: The binned summaries graph success proportion against perturb-target distance (relative to image width/height) and perturb box size (relative to image width/height) in the randomized attack experiment.

```
## List of 2
## $ mod : rowwise_df [15 x 4] (S3: rowwise_df/tbl_df/tbl/data.frame)
## ..$ model_name : Ord.factor w/ 5 levels "YOLOv3"<"SSD"<...: 1 1 1 2 2 2 3 3 3 4 ...
## ..$ loss_target: Ord.factor w/ 3 levels "Vanishing"<"Mislabeling"<...: 1 2 3 1 2 3 1 2 3 1 ...
## ..$ data : list<tibble[,15]> [1:15]
## ..$ mod :List of 15
## ..- attr(*, "groups")= tibble [15 x 3] (S3: tbl_df/tbl/data.frame)
## $ tidied: gropd_df [45 x 20] (S3: grouped_df/tbl_df/tbl/data.frame)
## ..$ model_name : Ord.factor w/ 5 levels "YOLOv3"<"SSD"<...: 1 1 1 1 1 1 1 1 1 2 ...
## ..$ loss_target : Ord.factor w/ 3 levels "Vanishing"<"Mislabeling"<...: 1 1 1 2 2 2 3 3 3 1 ...
## ..$ term : chr [1:45] "bbox_dist" "bbox_size_perturb" "bbox_dist:(bbox_size_perturb)" "bbox_size_perturb:(bbox_dist)" "bbox_dist:(bbox_size_perturb):bbox_size_perturb" ...
## ..$ variable : chr [1:45] "bbox_dist" "bbox_size_perturb" "bbox_dist:bbox_size_perturb" "bbox_size_perturb:bbox_dist" "bbox_dist:bbox_size_perturb:bbox_size_perturb" ...
## ..$ var_label : Named chr [1:45] "bbox_dist" "bbox_size_perturb" "bbox_dist * bbox_size_perturb" "bbox_size_perturb * bbox_dist" "bbox_dist * bbox_size_perturb * bbox_size_perturb" ...
## ..- attr(*, "names")= chr [1:45] "bbox_dist" "bbox_size_perturb" "bbox_dist:bbox_size_perturb" "bbox_size_perturb:bbox_dist" "bbox_dist:bbox_size_perturb:bbox_size_perturb" ...
## ..$ var_class : Named chr [1:45] "numeric" "numeric" NA "numeric" ...
## ..- attr(*, "names")= chr [1:45] "bbox_dist" "bbox_size_perturb" "" "bbox_dist" ...
## ..$ var_type : chr [1:45] "continuous" "continuous" "interaction" "continuous" ...
## ..$ var_nlevels : int [1:45] NA NA NA NA NA NA NA NA NA NA ...
## ..$ contrasts : chr [1:45] NA NA NA NA ...
```

```
## ..$ contrasts_type: chr [1:45] NA NA NA NA ...
## ..$ reference_row : logi [1:45] NA NA NA NA NA NA ...
## ..$ label : Named chr [1:45] "bbox_dist" "bbox_size_perturb" "bbox_dist * bbox_size_perturb"
## .. ..- attr(*, "names")= chr [1:45] "bbox_dist" "bbox_size_perturb" "bbox_dist:bbox_size_perturb"
## ..$ n_obs : Named num [1:45] 4000 4000 4000 4000 4000 4000 4000 4000 4000 4000 ...
## .. ..- attr(*, "names")= chr [1:45] "bbox_dist" "bbox_size_perturb" "bbox_dist:bbox_size_perturb"
## ..$ n_event : Named num [1:45] 1062 1062 1062 934 934 ...
## .. ..- attr(*, "names")= chr [1:45] "bbox_dist" "bbox_size_perturb" "bbox_dist:bbox_size_perturb"
## ..$ estimate : num [1:45] -8.54 26.83 -79.93 -8.47 10.99 ...
## ..$ std.error : num [1:45] 0.694 1.719 8.924 0.615 0.956 ...
## ..$ statistic : num [1:45] -12.29 15.61 -8.96 -13.78 11.5 ...
## ..$ p.value : num [1:45] 1.01e-34 6.26e-55 3.34e-19 3.45e-43 1.32e-30 ...
## ..$ conf.low : num [1:45] -9.93 23.55 -97.84 -9.71 9.17 ...
## ..$ conf.high : num [1:45] -7.21 30.29 -62.85 -7.3 12.92 ...
## ..- attr(*, "groups")= tibble [15 x 3] (S3: tbl_df/tbl/data.frame)
## .. ..- attr(*, ".drop")= logi TRUE
```

```
reg_est <- reg_res$tidied
```

```
ext_sig(reg_est, "neg", "bbox_dist")
```

```
## -----bbox_dist-----
## Total 15 predictors:
## 15 (100%) significant;
## 15 (100%) neg

## # A tibble: 15 x 9
## # Groups:   model_name, loss_target [15]
##   model_name loss_target term estimate std.error statistic p.value conf.low
##   <ord>      <ord>      <chr>   <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
## 1 YOLOv3     Vanishing  bbox~   -8.54    0.694    -12.3     0    -9.93
## 2 YOLOv3     Mislabeling bbox~   -8.47    0.615    -13.8     0    -9.71
## 3 YOLOv3     Untargeted  bbox~  -15.9    1.37     -11.6     0   -18.6
## 4 SSD        Vanishing  bbox~  -18.4    1.16     -15.9     0  -20.8
## 5 SSD        Mislabeling bbox~  -19.7    1.31     -15.0     0  -22.3
## 6 SSD        Untargeted  bbox~  -21.7    1.54     -14.1     0  -24.9
## 7 RetinaNet  Vanishing  bbox~  -35.3    3.25     -10.9     0  -41.9
## 8 RetinaNet  Mislabeling bbox~  -49.8    6.49      -7.68     0  -63.3
## 9 RetinaNet  Untargeted  bbox~  -13.9    1.41      -9.84     0  -16.8
## 10 Faster R-CNN Vanishing  bbox~  -21.0    3.20      -6.56     0  -27.7
## 11 Faster R-CNN Mislabeling bbox~  -17.8    3.24      -5.51     0  -24.7
## 12 Faster R-CNN Untargeted  bbox~  -19.1    1.79     -10.7     0  -22.7
## 13 Cascade R-CNN Vanishing  bbox~  -32.5    4.07      -7.99     0  -41.0
## 14 Cascade R-CNN Mislabeling bbox~  -27.7    4.73      -5.86     0  -37.8
## 15 Cascade R-CNN Untargeted  bbox~  -22.5    2.47      -9.12     0  -27.6
## # i 1 more variable: conf.high <dbl>
```

```
ext_sig(reg_est, "pos", "bbox_size_perturb")
```

```
## -----bbox_size_perturb-----
## Total 15 predictors:
## 11 (73%) significant;
## 11 (73%) pos

## # A tibble: 11 x 9
## # Groups:   model_name, loss_target [11]
```

```
##   model_name   loss_target term   estimate std.error statistic p.value conf.low
##   <ord>        <ord>      <chr>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
## 1 YOLOv3       Vanishing  bbox~    26.8     1.72     15.6     0        23.6
## 2 YOLOv3       Mislabeling bbox~    11.0     0.956    11.5     0         9.17
## 3 SSD          Vanishing  bbox~     7.27    0.813     8.95     0         5.73
## 4 SSD          Mislabeling bbox~     3.38    0.612     5.53     0         2.22
## 5 SSD          Untargeted bbox~     1.39    0.545     2.55    0.011     0.336
## 6 RetinaNet    Vanishing  bbox~     2.32    0.695     3.33    0.001     0.993
## 7 RetinaNet    Untargeted bbox~     2.99    0.539     5.54     0         1.94
## 8 Faster R-CNN Vanishing  bbox~     6.10    1.23      4.96     0         3.75
## 9 Cascade R-CNN Vanishing  bbox~     7.51    0.966     7.78     0         5.71
##10 Cascade R-CNN Mislabeling bbox~     4.90    0.797     6.15     0         3.35
##11 Cascade R-CNN Untargeted bbox~     2.11    0.648     3.26    0.001     0.833
## # i 1 more variable: conf.high <dbl>
```

```
ext_sig(reg_est, "both", "bbox_dist:(bbox_size_perturb)")
```

```
## -----bbox_dist:(bbox_size_perturb)-----
```

```
## Total 15 predictors:
## 10 (67%) significant;
## 10 (67%) both
```

```
## # A tibble: 10 x 9
```

```
## # Groups:   model_name, loss_target [10]
```

```
##   model_name   loss_target term   estimate std.error statistic p.value conf.low
##   <ord>        <ord>      <chr>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
## 1 YOLOv3       Vanishing  bbox~   -79.9     8.92     -8.96     0      -97.8
## 2 YOLOv3       Mislabeling bbox~   -24.1     5.92     -4.08     0     -36.0
## 3 YOLOv3       Untargeted bbox~    39.5     6.52      6.06     0      26.7
## 4 SSD          Mislabeling bbox~    24.0     6.04      3.97     0      12.0
## 5 SSD          Untargeted bbox~    34.2     6.42      5.32     0      21.4
## 6 RetinaNet    Vanishing  bbox~    47.0    11.2      4.19     0      24.3
## 7 RetinaNet    Untargeted bbox~    28.1     5.11      5.49     0      18.1
## 8 Faster R-CNN Vanishing  bbox~   -83.5    28.5     -2.93    0.003  -144.
## 9 Faster R-CNN Untargeted bbox~    61.5     6.97      8.82     0      48.4
##10 Cascade R-CNN Vanishing  bbox~  -106.    31.1     -3.42    0.001  -172.
## # i 1 more variable: conf.high <dbl>
```

```
print_statistics(reg_est, table_caption(pred_name, main_pt))
```

Table 5: We run a logistic model regressing success against perturb-target distance (relative to image width/height) and perturb box size (relative to image width/height) in the randomized attack experiment. Larger perturb objects significantly increase success rates for all models and attacks, except for mislabeling attack on Faster R-CNN, after controlling for perturb-target distances. Shorter perturb-target distances significantly increase success rates for all models and attacks, after controlling for perturb object sizes. Table headers are explained in Appendix ??.

Group		Regression						
Attack	term	sig	estimate	std.error	statistic	p.value	conf.low	conf.high
<b>YOLOv3</b>	Vanishing							
	distance	*	-8.536	0.694	-12.292	0.000	-9.929	-7.207
	size	*	26.831	1.719	15.610	0.000	23.555	30.294
	distance * size	*	-79.933	8.924	-8.957	0.000	-97.839	-62.847

Mislabeling	distance	*	-8.473	0.615	-13.778	0.000	-9.707	-7.297
	size	*	10.991	0.956	11.500	0.000	9.169	12.915
	distance * size	*	-24.117	5.917	-4.076	0.000	-35.972	-12.770
Untargeted	distance	*	-15.869	1.366	-11.614	0.000	-18.640	-13.284
	size		0.308	0.704	0.437	0.662	-1.087	1.678
	distance * size	*	39.532	6.522	6.061	0.000	26.743	52.347
<b>SSD</b>								
Vanishing	distance	*	-18.433	1.159	-15.903	0.000	-20.766	-16.222
	size	*	7.274	0.813	8.948	0.000	5.728	8.915
	distance * size		7.663	6.391	1.199	0.231	-5.139	19.931
Mislabeling	distance	*	-19.702	1.311	-15.023	0.000	-22.349	-17.208
	size	*	3.384	0.612	5.531	0.000	2.217	4.617
	distance * size	*	23.987	6.040	3.971	0.000	11.954	35.660
Untargeted	distance	*	-21.725	1.544	-14.069	0.000	-24.852	-18.799
	size	*	1.389	0.545	2.547	0.011	0.336	2.478
	distance * size	*	34.171	6.423	5.320	0.000	21.425	46.643
<b>RetinaNet</b>								
Vanishing	distance	*	-35.303	3.249	-10.864	0.000	-41.932	-29.191
	size	*	2.317	0.695	3.334	0.001	0.993	3.717
	distance * size	*	46.975	11.215	4.189	0.000	24.285	68.263
Mislabeling	distance	*	-49.847	6.486	-7.685	0.000	-63.277	-37.849
	size		1.056	1.187	0.889	0.374	-1.244	3.427
	distance * size		37.912	25.512	1.486	0.137	-15.784	84.709
Untargeted	distance	*	-13.895	1.412	-9.843	0.000	-16.788	-11.254
	size	*	2.989	0.539	5.544	0.000	1.938	4.054
	distance * size	*	28.072	5.111	5.493	0.000	18.127	38.241
<b>Faster R-CNN</b>								
Vanishing	distance	*	-21.030	3.204	-6.564	0.000	-27.739	-15.185
	size	*	6.096	1.228	4.962	0.000	3.747	8.571
	distance * size	*	-83.474	28.510	-2.928	0.003	-144.255	-31.915
Mislabeling	distance	*	-17.846	3.240	-5.507	0.000	-24.720	-12.034
	size		1.205	1.719	0.701	0.483	-2.408	4.397
	distance * size		-54.135	39.695	-1.364	0.173	-142.163	14.635
Untargeted	distance	*	-19.078	1.789	-10.665	0.000	-22.746	-15.729
	size		-0.274	0.719	-0.381	0.703	-1.711	1.113
	distance * size	*	61.468	6.966	8.824	0.000	48.369	75.700
<b>Cascade R-CNN</b>								
Vanishing	distance	*	-32.490	4.066	-7.991	0.000	-40.976	-25.029
	size	*	7.513	0.966	7.779	0.000	5.711	9.508
	distance * size	*	-106.218	31.092	-3.416	0.001	-172.083	-49.911
Mislabeling	distance	*	-27.708	4.732	-5.856	0.000	-37.836	-19.260
	size	*	4.898	0.797	6.146	0.000	3.354	6.485

			distance * size	-49.344	27.328	-1.806	0.071	-107.414	-0.192
Untargeted	distance	*	-22.497	2.467	-9.120	0.000	-27.587	-17.915	
	size	*	2.113	0.648	3.258	0.001	0.833	3.381	
	distance * size		5.873	11.482	0.512	0.609	-18.022	27.276	

```
reg_mod <- reg_res$mod

newdata <- expand_grid(
  bbox_dist = linear_space(data$bbox_dist),
  bbox_size_perturb = linear_space(data$bbox_size_perturb)
) |>
  glimpse()

## Rows: 10,000
## Columns: 2
## $ bbox_dist      <dbl> 1.180172e-05, 1.180172e-05, 1.180172e-05, 1.180172e-~
## $ bbox_size_perturb <dbl> 2.671581e-05, 8.839428e-03, 1.765214e-02, 2.646485e-~

reg_pred <- reg_mod |>
  summarize(augment(mod, newdata = newdata, type.predict = "response")) |>
  rename(success = .fitted) |>
  glimpse()

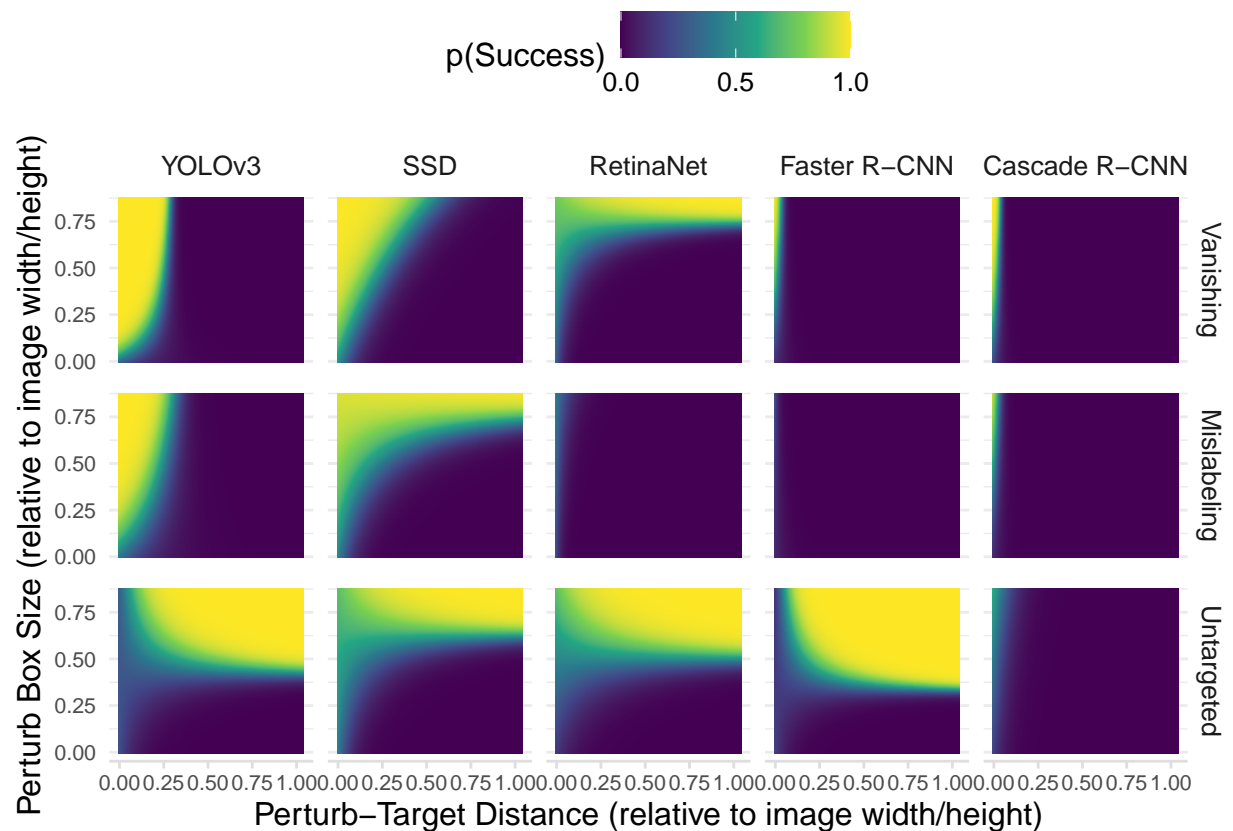
## Warning: Returning more (or less) than 1 row per `summarise()` group was deprecated in
## dplyr 1.1.0.
## i Please use `reframe()` instead.
## i When switching from `summarise()` to `reframe()`, remember that `reframe()`
## always returns an ungrouped data frame and adjust accordingly.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.

## `summarise()` has grouped output by 'model_name', 'loss_target'. You can
## override using the `.groups` argument.

## Rows: 150,000
## Columns: 5
## Groups: model_name, loss_target [15]
## $ model_name      <ord> YOLOv3, YOLOv3, YOLOv3, YOLOv3, YOLOv3, YOLOv3, YOLO~
## $ loss_target     <ord> Vanishing, Vanishing, Vanishing, Vanishing, Vanishin~
## $ bbox_dist       <dbl> 1.180172e-05, 1.180172e-05, 1.180172e-05, 1.180172e-~
## $ bbox_size_perturb <dbl> 2.671581e-05, 8.839428e-03, 1.765214e-02, 2.646485e-~
## $ success         <dbl> 0.3376685, 0.3923941, 0.4499635, 0.5089050, 0.567600~

g <- reg_pred |> ggplot(aes(bbox_dist, bbox_size_perturb, fill = success)) +
  geom_raster(interpolate = TRUE)

graph_dist_size(g)
```



```
# get success rate on ground truth sampled images
gt_success_data <- bbox_raw_data |>
  filter(bbox_type == "ground_truth") |>
  # loss_target is not relevant
  count(model_name, bbox_class, bbox_res_eval) |>
  # get success probability
  # https://stackoverflow.com/a/37448040/19655086
  as_tibble() |>
  pivot_wider(names_from = "bbox_res_eval", values_from = n) |>
  # not every class has tp and fn
  replace_na(list(tp = 0, fn = 0)) |>
  mutate(gt_p_success = tp / (fn + tp)) |>
  # some 0/0
  drop_na(gt_p_success) |>
  select(model_name, bbox_class, gt_p_success) |>
  glimpse()
```

```
## Rows: 378
## Columns: 3
## $ model_name   <ord> YOLOv3, YOLOv3, YOLOv3, YOLOv3, YOLOv3, YOLOv3, YOLOv3, Y-
## $ bbox_class   <chr> "airplane", "apple", "backpack", "banana", "baseball bat"~
## $ gt_p_success <dbl> 1.0000000, 0.7196110, 0.2938005, 0.1124452, 0.5025961, 0.~
```

```
# by model_name, bbox_class
# some classes are not evaluated
gt_success_data <- bbox_conf_data |>
  inner_join(gt_success_data) |>
  glimpse()
```

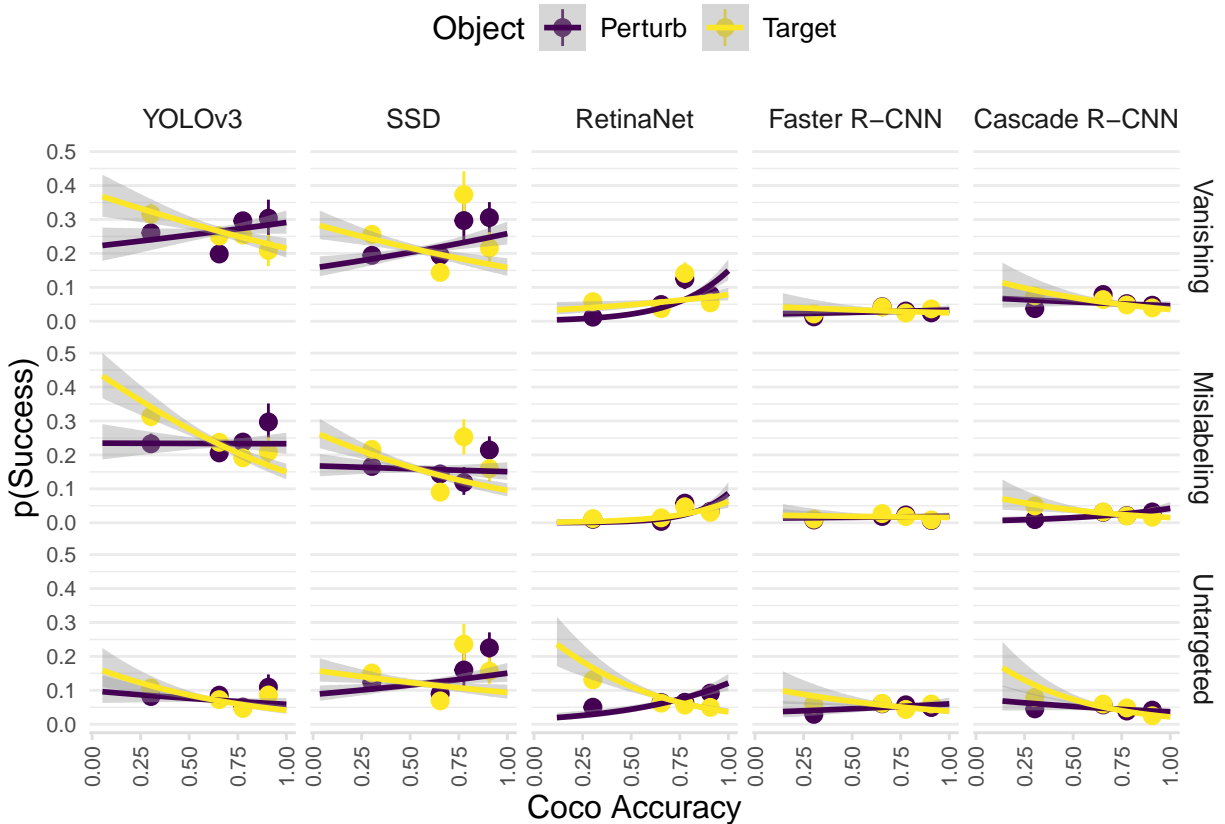


```

## Joining with `by = join_by(bbox_class, model_name)`

## Rows: 120,000
## Columns: 43
## $ fname <chr> "/Users/zbli/Documents/Documents - ZhaoBin"~
## $ sample_id <chr> "65ed3a88141a475067f32706", "65ed3a88141a47~
## $ sample_path <chr> "/projects/f_ps848_1/zhaobin/adversarial/co~
## $ sample_width <int> 640, 640, 500, 640, 480, 640, 640, 640, 640~
## $ sample_height <int> 480, 427, 332, 425, 640, 480, 480, 480, 640~
## $ sample_mislabel_class <chr> "horse", "truck", "surfboard", "horse", "ca~
## $ sample_mislabel_proba <dbl> 6.615031e-05, 4.219168e-02, 4.392489e-05, 1~
## $ sample_attack <lg1> TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, T~
## $ sample_vanish <lg1> FALSE, FALSE, TRUE, FALSE, FALSE, FALSE, FA~
## $ sample_mislabel_intended <lg1> FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, F~
## $ sample_success <lg1> FALSE, FALSE, TRUE, FALSE, FALSE, FALSE, FA~
## $ sample_mislabel <lg1> FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, F~
## $ bbox_id <chr> "65ed3aa3141a475067f3ca3e", "65ed3aa3141a47~
## $ bbox_class <chr> "clock", "car", "person", "person", "donut"~
## $ bbox_xywhn <list<double>> <0.32723613, 0.26601949, 0.0435188~
## $ bbox_conf <dbl> 0.9305881, 0.3433506, 0.9882318, 0.9988949,~
## $ bbox_res_eval <chr> "tp", "tp", "tp", "tp", "tp", "tp", "tp", "~
## $ bbox_iou_eval <dbl> 0.8860679, 0.7609860, 0.9454082, 0.9299325,~
## $ bbox_res_pgd_eval <chr> "tp", "tp", "fn", "tp", "tp", "tp", "tp", "~
## $ bbox_iou_pgd_eval <dbl> 1.0000000, 1.0000000, NA, 0.9999969, 1.0000~
## $ bbox_res_pgd_mislabel_eval <chr> NA, NA, "fn", NA, NA, NA, NA, NA, NA, N~
## $ bbox_iou_pgd_mislabel_eval <dbl> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ bbox_type <chr> "predictions", "predictions", "predictions"~
## $ bbox_mislabel <lg1> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ num_iteration <dbl> 200, 200, 200, 200, 200, 200, 200, 200, 200~
## $ max_norm <dbl> 0.05, 0.05, 0.05, 0.05, 0.05, 0.05, 0.05, 0~
## $ model_name <ord> Cascade R-CNN, Cascade R-CNN, Cascade R-CNN~
## $ loss_target <ord> Mislabeling, Mislabeling, Mislabeling, Misl~
## $ attack_bbox <chr> "predictions", "predictions", "predictions"~
## $ perturb_fun <chr> "perturb_inside", "perturb_inside", "pertur~
## $ sample_count <dbl> 247, 247, 247, 247, 247, 247, 247, 247, 247~
## $ attack_count <dbl> 200, 200, 200, 200, 200, 200, 200, 200, 200~
## $ success_count <dbl> 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8, 8~
## $ vanish_count <dbl> 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2~
## $ mislabel_count <dbl> 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6~
## $ mislabel_intended_count <dbl> 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6~
## $ target_max_conf <lg1> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ perturb_min_size <lg1> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ bbox_max_dist <lg1> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ target_or_perturb <ord> Target, Target, Target, Target, Target, Tar~
## $ target_or_perturb_boolean <lg1> TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, T~
## $ success <dbl> 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
## $ gt_p_success <dbl> 0.9090909, 0.7741935, 0.8199719, 0.8199719,~
gt_success_data |>
  graph_attr(gt_p_success, "COCO Accuracy")

```



```
pred_name <- "mean COCO accuracy for the target class"
main_pt <- "the results are mixed after controlling for target class confidence"

cap <- graph_caption(pred_name, glue("Although higher {pred_name} seem to decrease success rates, {main_pt}"))

gt_success_graph <- gt_success_data |> filter(target_or_perturb == "Target")
gt_success_graph |>
  graph_attr(gt_p_success, pred_name)

model <- partial(glm_model, predictor = "gt_p_success * bbox_conf")
data <- gt_success_graph

reg_est <- get_tidied_reg(model, data)

## Warning: Returning more (or less) than 1 row per `summarise()` group was deprecated in
## dplyr 1.1.0.
## i Please use `reframe()` instead.
## i When switching from `summarise()` to `reframe()`, remember that `reframe()`
## always returns an ungrouped data frame and adjust accordingly.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.

## `summarise()` has grouped output by 'model_name', 'loss_target'. You can
## override using the `.groups` argument.
# there are both significantly positive and negative gt_p_success,
# and the interaction term is relatively large
ext_sig(reg_est, "neg", "gt_p_success")
```

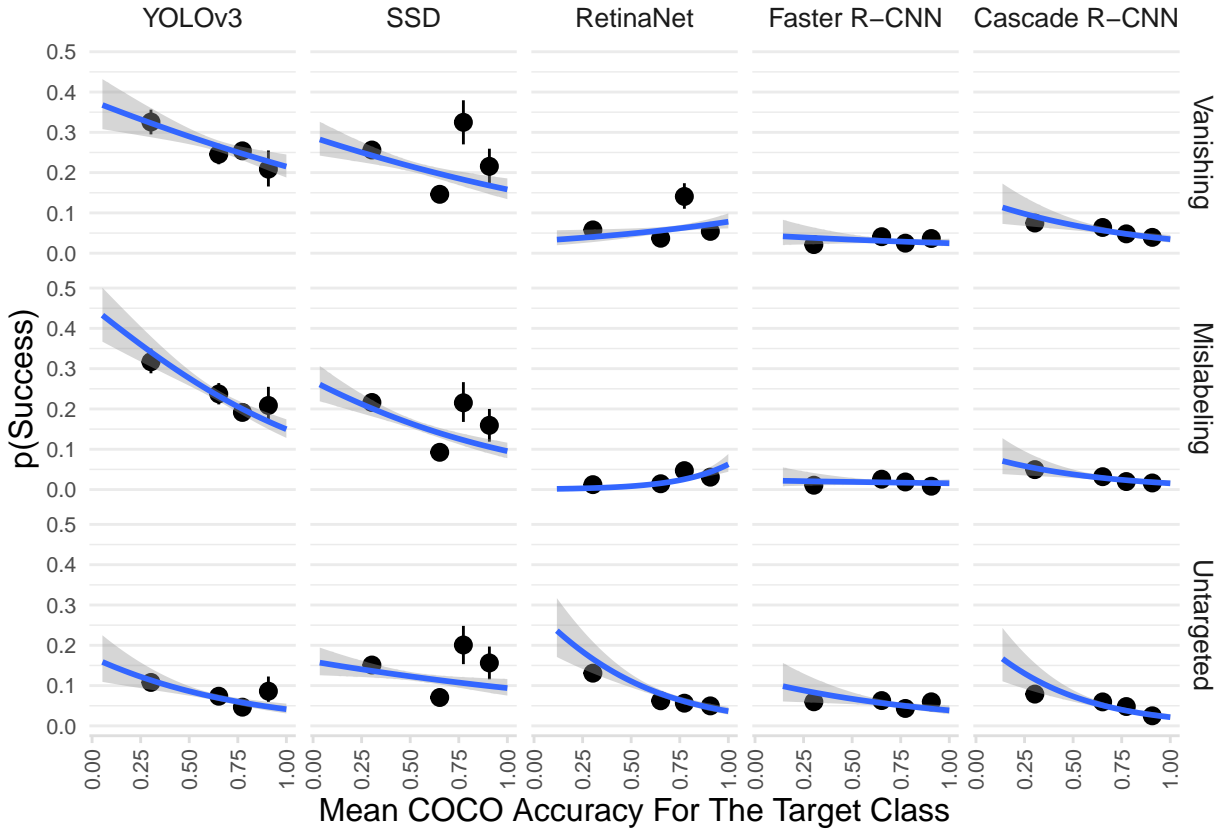


Figure 4: Although higher mean COCO accuracy for the target class seem to decrease success rates, the results are mixed after controlling for target class confidence (Table 6) even with 0.05 max-norm: The binned summaries and regression trendlines graph success proportion against mean COCO accuracy for the target class in the randomized attack experiment. Bins are split into quantiles. Errors are 95% confidence intervals

```
## -----gt_p_success-----
## Total 15 predictors:
## 8 (53%) significant;
## 2 (13%) neg

## # A tibble: 2 x 9
## # Groups:   model_name, loss_target [2]
##   model_name    loss_target term      estimate std.error statistic p.value conf.low
##   <ord>         <ord>      <chr>      <dbl>      <dbl>      <dbl>  <dbl>   <dbl>
## 1 Cascade R-CNN Vanishing  gt_p_~    -4.25      1.49      -2.85   0.004   -7.16
## 2 Cascade R-CNN Mislabeling gt_p_~    -4.57      1.81      -2.53   0.011   -8.08
## # i 1 more variable: conf.high <dbl>

ext_sig(reg_est, "pos", "gt_p_success")

## -----gt_p_success-----
## Total 15 predictors:
## 8 (53%) significant;
## 6 (40%) pos

## # A tibble: 6 x 9
## # Groups:   model_name, loss_target [6]
##   model_name loss_target term      estimate std.error statistic p.value conf.low
```

```
##   <ord>      <ord>      <chr>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
## 1 SSD       Vanishing   gt_p_suc~  3.77      0.582      6.48      0          2.64
## 2 SSD       Mislabeling gt_p_suc~  4.38      0.63       6.95      0          3.15
## 3 SSD       Untargeted  gt_p_suc~  3.38      0.681      4.96      0          2.05
## 4 RetinaNet Vanishing   gt_p_suc~  3.27      1.39       2.35      0.019      0.576
## 5 RetinaNet Mislabeling gt_p_suc~ 11.0      2.73       4.02      0          5.68
## 6 RetinaNet Untargeted  gt_p_suc~  3.55      1.29       2.75      0.006      1.03
## # i 1 more variable: conf.high <dbl>

ext_sig(reg_est, "both", "gt_p_success:bbbox_conf")

## -----gt_p_success:bbbox_conf-----
## Total 15 predictors:
## 10 (67%) significant;
## 10 (67%) both

## # A tibble: 10 x 9
## # Groups:   model_name, loss_target [10]
##   model_name  loss_target term  estimate std.error statistic p.value conf.low
##   <ord>      <ord>      <chr>   <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
## 1 YOLOv3     Vanishing   gt_p~   -2.05    1.01     -2.03    0.042    -4.03
## 2 YOLOv3     Mislabeling gt_p~   -3.48    1.06     -3.27    0.001    -5.57
## 3 YOLOv3     Untargeted  gt_p~   -4.86    1.91     -2.54    0.011    -8.61
## 4 SSD        Vanishing   gt_p~   -6.66    0.854    -7.79    0        -8.34
## 5 SSD        Mislabeling gt_p~   -8.65    0.976    -8.86    0        -10.6
## 6 SSD        Untargeted  gt_p~   -6.06    1.11     -5.48    0        -8.24
## 7 RetinaNet  Mislabeling gt_p~  -11.7    5.71     -2.05    0.04     -22.6
## 8 RetinaNet  Untargeted  gt_p~   -9.35    2.76     -3.39    0.001    -14.8
## 9 Cascade R-CNN Vanishing   gt_p~    4.33    1.96      2.21    0.027     0.483
##10 Cascade R-CNN Mislabeling gt_p~    5.32    2.64      2.02    0.044     0.152
## # i 1 more variable: conf.high <dbl>

print_statistics(reg_est, table_caption(
  glue("{pred_name}, with target confidence as covariate,"),
  glue("{main_pt} and the relatively large interaction terms make interpretation challenging")
))
```

Table 6: We run a logistic model regressing success against mean COCO accuracy for the target class, with target confidence as covariate, in the randomized attack experiment. The results are mixed after controlling for target class confidence and the relatively large interaction terms make interpretation challenging. Table headers are explained in Appendix ??.

Group		Regression						
Attack	term	sig	estimate	std.error	statistic	p.value	conf.low	conf.high
<b>YOLOv3</b>								
Vanishing	accuracy		0.842	0.747	1.127	0.260	-0.619	2.313
	confidence		0.368	0.671	0.548	0.584	-0.945	1.688
	accuracy * confidence	*	-2.046	1.007	-2.031	0.042	-4.026	-0.076
Mislabeling	accuracy		1.231	0.754	1.631	0.103	-0.247	2.712
	confidence		-0.139	0.700	-0.198	0.843	-1.514	1.234
	accuracy * confidence	*	-3.481	1.065	-3.270	0.001	-5.571	-1.396
Untargeted	accuracy		1.941	1.117	1.737	0.082	-0.240	4.143

		confidence		-1.715	1.230	-1.394	0.163	-4.155	0.671
		accuracy * confidence	*	-4.861	1.913	-2.541	0.011	-8.612	-1.112
<b>SSD</b>									
Vanishing	accuracy		*	3.774	0.582	6.485	0.000	2.640	4.923
	confidence		*	2.184	0.491	4.451	0.000	1.226	3.150
	accuracy * confidence		*	-6.655	0.854	-7.789	0.000	-8.340	-4.990
Mislabeling	accuracy		*	4.376	0.630	6.950	0.000	3.148	5.618
	confidence		*	2.449	0.538	4.550	0.000	1.395	3.506
	accuracy * confidence		*	-8.650	0.976	-8.864	0.000	-10.573	-6.746
Untargeted	accuracy		*	3.376	0.681	4.955	0.000	2.047	4.720
	confidence			0.423	0.626	0.677	0.499	-0.809	1.646
	accuracy * confidence		*	-6.063	1.106	-5.480	0.000	-8.239	-3.902
<b>RetinaNet</b>									
Vanishing	accuracy		*	3.267	1.389	2.353	0.019	0.576	6.018
	confidence			-0.776	2.077	-0.374	0.709	-4.879	3.260
	accuracy * confidence			-2.512	2.651	-0.948	0.343	-7.702	2.686
Mislabeling	accuracy		*	10.978	2.731	4.020	0.000	5.683	16.358
	confidence			3.473	4.602	0.755	0.450	-5.826	12.146
	accuracy * confidence		*	-11.692	5.707	-2.049	0.040	-22.608	-0.344
Untargeted	accuracy		*	3.553	1.292	2.751	0.006	1.029	6.093
	confidence			0.863	1.920	0.449	0.653	-2.964	4.566
	accuracy * confidence		*	-9.351	2.760	-3.388	0.001	-14.760	-3.935
<b>Faster R-CNN</b>									
Vanishing	accuracy			-1.752	1.802	-0.973	0.331	-5.202	1.874
	confidence		*	-6.201	2.110	-2.939	0.003	-10.372	-2.093
	accuracy * confidence			3.626	2.762	1.313	0.189	-1.797	9.030
Mislabeling	accuracy			2.740	2.469	1.110	0.267	-1.989	7.689
	confidence			-3.313	3.126	-1.060	0.289	-9.642	2.613
	accuracy * confidence			-2.724	4.126	-0.660	0.509	-10.668	5.473
Untargeted	accuracy			1.841	1.415	1.301	0.193	-0.897	4.655
	confidence			-2.543	1.607	-1.583	0.114	-5.733	0.572
	accuracy * confidence			-2.728	2.162	-1.262	0.207	-6.949	1.529
<b>Cascade R-CNN</b>									
Vanishing	accuracy		*	-4.247	1.491	-2.848	0.004	-7.156	-1.298
	confidence		*	-4.563	1.413	-3.229	0.001	-7.328	-1.779
	accuracy * confidence		*	4.330	1.956	2.214	0.027	0.483	8.158
Mislabeling	accuracy		*	-4.568	1.806	-2.530	0.011	-8.081	-0.985
	confidence		*	-6.823	1.939	-3.519	0.000	-10.663	-3.046
	accuracy * confidence		*	5.322	2.638	2.017	0.044	0.152	10.503
Untargeted	accuracy			-0.017	1.423	-0.012	0.990	-2.791	2.794
	confidence			-1.750	1.449	-1.207	0.227	-4.607	1.083

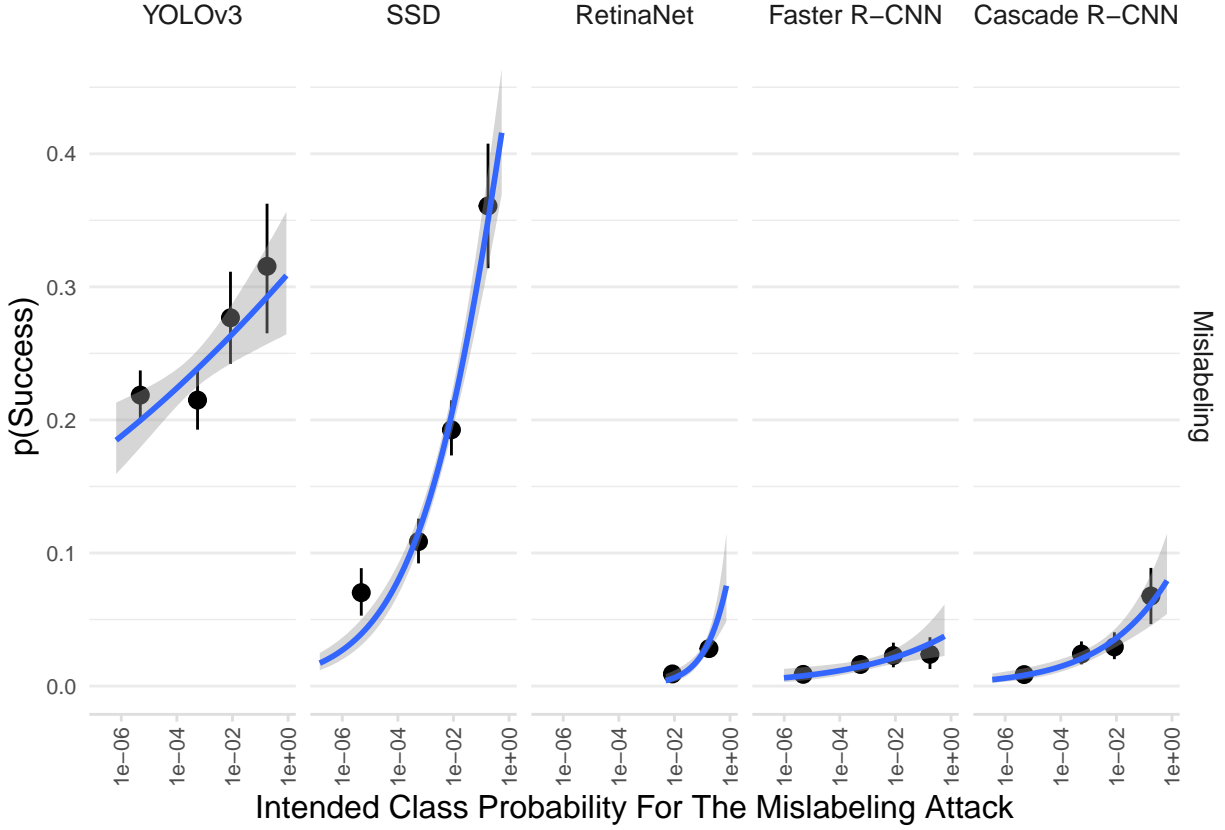


Figure 5: Although intended class probability seem to increase success rates for the mislabeling attack, it does not predict success rates after controlling for target class confidence, except for RetinaNet (Table 7) even with 0.05 max-norm: The binned summaries and regression trendlines graph success proportion against intended class probability in the randomized attack experiment. Bins are split into quantiles. Errors are 95% confidence intervals

accuracy * confidence	-2.732	2.037	-1.341	0.180	-6.726	1.265
-----------------------	--------	-------	--------	-------	--------	-------

```
# restrict to mislabeling
bbox_proba_graph <- bbox_conf_data |>
  filter(loss_target == "Mislabeling" & target_or_perturb == "Target")

# check is not logit
stopifnot(max(bbox_proba_graph$sample_mislabel_proba) <= 1 && min(bbox_proba_graph$sample_mislabel_proba) >= 0)

pred_name <- "intended class probability"
att_name <- "for the mislabeling attack"

main_pt <- glue("does not predict success rates after controlling for target class confidence, except for RetinaNet")
cap <- graph_caption(pred_name, glue("Although {pred_name} seem to increase success rates {att_name}, it does not predict success rates after controlling for target class confidence, except for RetinaNet"))

g <- bbox_proba_graph |>
  graph_attr(sample_mislabel_proba, glue("{pred_name} {att_name}"), scale_x_log10())

model <- partial(glm_model, predictor = "log(sample_mislabel_proba) * bbox_conf")
data <- bbox_proba_graph
```

```

reg_est <- get_tidied_reg(model, data)

## Warning: Returning more (or less) than 1 row per `summarise()` group was deprecated in
## dplyr 1.1.0.
## i Please use `reframe()` instead.
## i When switching from `summarise()` to `reframe()`, remember that `reframe()`
## always returns an ungrouped data frame and adjust accordingly.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.

## `summarise()` has grouped output by 'model_name', 'loss_target'. You can
## override using the `.groups` argument.

ext_sig(reg_est, "pos", "log(sample_mislabel_proba)")

## -----log(sample_mislabel_proba)-----
## Total 5 predictors:
## 3 (60%) significant;
## 2 (40%) pos

## # A tibble: 2 x 9
## # Groups:   model_name, loss_target [2]
##   model_name loss_target term          estimate std.error statistic p.value conf.low
##   <ord>      <ord>      <chr>          <dbl>     <dbl>     <dbl>   <dbl>   <dbl>
## 1 SSD       Mislabeling log(samp~    0.196     0.055      3.57    0       0.089
## 2 RetinaNet Mislabeling log(samp~    1.12     0.373      2.99    0.003    0.374
## # i 1 more variable: conf.high <dbl>

ext_sig(reg_est, "both", "log(sample_mislabel_proba):bbox_conf")

## -----log(sample_mislabel_proba):bbox_conf-----
## Total 5 predictors:
## 1 (20%) significant;
## 1 (20%) both

## # A tibble: 1 x 9
## # Groups:   model_name, loss_target [1]
##   model_name loss_target term          estimate std.error statistic p.value conf.low
##   <ord>      <ord>      <chr>          <dbl>     <dbl>     <dbl>   <dbl>   <dbl>
## 1 YOLOv3     Mislabeling log(samp~    0.317     0.062      5.14    0       0.196
## # i 1 more variable: conf.high <dbl>

print_statistics(reg_est, table_caption(glue("log({pred_name}) {att_name}, with predicted class's confi


```

Table 7: We run a logistic model regressing success against log(intended class probability) for the mislabeling attack, with predicted class's confidence as covariate, in the randomized attack experiment. Intended class probability does not predict success rates after controlling for target class confidence, except for RetinaNet. Table headers are explained in Appendix ??.

Group		Regression							
Model	term	sig	estimate	std.error	statistic	p.value	conf.low	conf.high	
Mislabeling	YOLOv3								
		log(probability)	*	-0.183	0.042	-4.344	0.000	-0.266	-0.101
		confidence		0.119	0.522	0.227	0.820	-0.904	1.143

	log(probability) * confidence	*	0.317	0.062	5.140	0.000	0.196	0.438
SSD	log(probability)	*	0.196	0.055	3.574	0.000	0.089	0.304
	confidence	*	-1.546	0.503	-3.071	0.002	-2.532	-0.558
	log(probability) * confidence		0.011	0.078	0.146	0.884	-0.141	0.166
RetinaNet	log(probability)	*	1.117	0.373	2.993	0.003	0.374	1.837
	confidence	*	-8.002	1.997	-4.006	0.000	-11.970	-4.136
	log(probability) * confidence		-1.384	0.757	-1.828	0.067	-2.822	0.145
Faster R-CNN	log(probability)		0.158	0.120	1.314	0.189	-0.080	0.393
	confidence	*	-7.667	1.544	-4.964	0.000	-10.765	-4.692
	log(probability) * confidence		-0.330	0.196	-1.684	0.092	-0.709	0.061
Cascade R-CNN	log(probability)		0.096	0.111	0.864	0.388	-0.123	0.313
	confidence	*	-2.499	1.024	-2.440	0.015	-4.493	-0.470
	log(probability) * confidence		0.020	0.153	0.133	0.894	-0.275	0.326

```
# bbox iou always based on predictions bbox like confidence
bbox_conf_data |>
  graph_attr(bbox_iou_eval, " IOU ")
```



```
# restrict to target bbox and untargeted attack only
pred_name <- "target iou for the untargeted attack"
main_pt <- glue("{pred_name} increases success rates on all models")

cap <- graph_caption(pred_name, main_pt, params$norm)
```






Figure 6: Target IOU for the untargeted attack increases success rates on all models even with 0.05 max-norm: The binned summaries and regression trendlines graph success proportion against target IOU for the untargeted attack in the randomized attack experiment. Bins are split into quantiles. Errors are 95% confidence intervals

```
bbox_iou_graph <- bbox_conf_data |> filter(target_or_perturb == "Target" & loss_target == "Untargeted")
bbox_iou_graph |>
  graph_attr(bbox_iou_eval, pred_name)
```

```
model <- partial(glm_model, predictor = "bbox_iou_eval")
data <- bbox_iou_graph
```

```
reg_est <- get_tidied_reg(model, data)
```

```
## `summarise()` has grouped output by 'model_name', 'loss_target'. You can
## override using the `.groups` argument.
```

```
ext_sig(reg_est, "neg")
```

```
## Total 5 predictors:
## 5 (100%) significant;
## 5 (100%) neg

## # A tibble: 5 x 9
## # Groups:   model_name, loss_target [5]
##   model_name loss_target term estimate std.error statistic p.value conf.low
##   <ord>      <ord>      <chr>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
```

```
## 1 YOLOv3      Untargeted  bbox_~   -3.19    0.351    -9.10     0    -3.88
## 2 SSD         Untargeted  bbox_~   -2.75    0.288    -9.54     0    -3.31
## 3 RetinaNet   Untargeted  bbox_~   -3.08    0.328    -9.40     0    -3.72
## 4 Faster R-CNN Untargeted  bbox_~   -2.02    0.374    -5.40     0    -2.74
## 5 Cascade R-CNN Untargeted  bbox_~   -2.90    0.364    -7.95     0    -3.61
## # i 1 more variable: conf.high <dbl>
```

```
print_statistics(reg_est, table_caption(pred_name, main_pt))
```

Table 8: We run a logistic model regressing success against target IOU for the untargeted attack in the randomized attack experiment. Target IOU for the untargeted attack increases success rates on all models. Table headers are explained in Appendix ??.

Group		Regression						
Model	term	sig	estimate	std.error	statistic	p.value	conf.low	conf.high
<b>Untargeted</b>								
YOLOv3	bbox_iou_eval	*	-3.194	0.351	-9.098	0	-3.878	-2.501
SSD	bbox_iou_eval	*	-2.747	0.288	-9.539	0	-3.309	-2.180
RetinaNet	bbox_iou_eval	*	-3.085	0.328	-9.402	0	-3.725	-2.438
Faster R-CNN	bbox_iou_eval	*	-2.020	0.374	-5.403	0	-2.745	-1.278
Cascade R-CNN	bbox_iou_eval	*	-2.895	0.364	-7.953	0	-3.606	-2.177